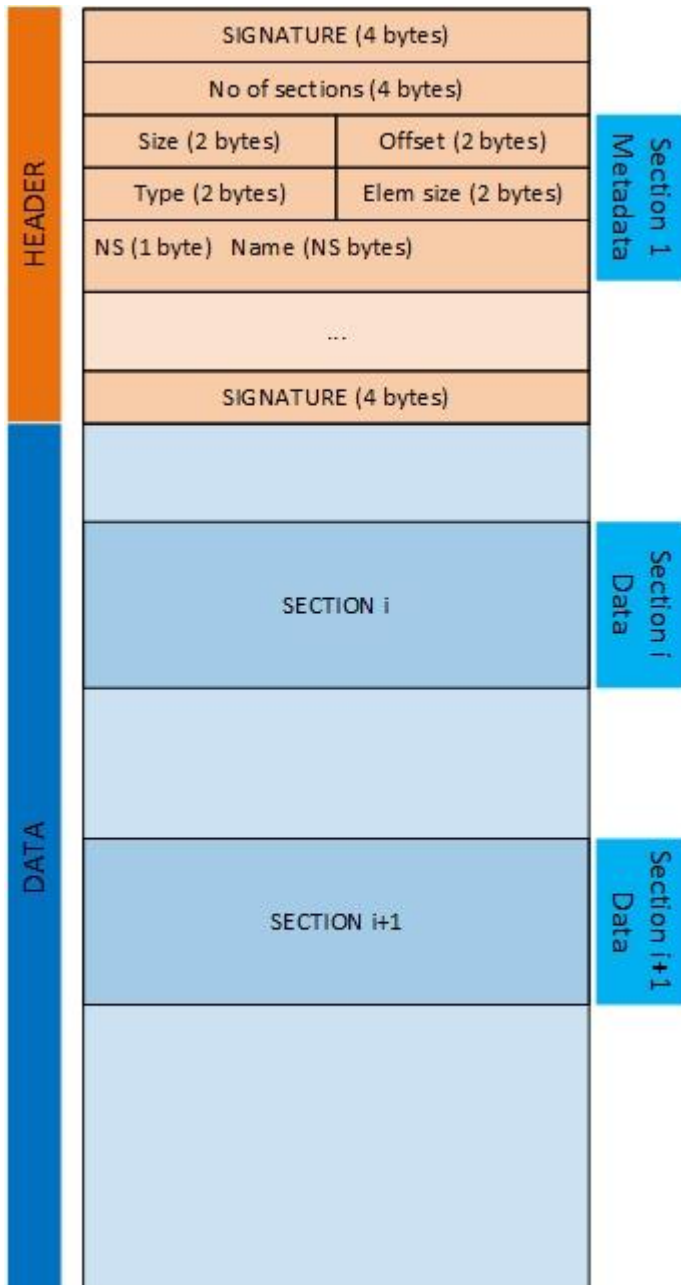


Lab ASSIGNMENT 1

You are given the following file format, which we will call from now on **SF** (i.e. "**section file**") **format**. A section file consists in two areas: a **header** and a **data area**. The header area describes the way the data area must be read. The header area also contains the file's signature, both at its beginning and end, like illustrated in the figure below. The signature of the section files is "0xB612B612".



The data area consists in more **sections**, not necessarily placed one after another, i.e. there could be holes between sections. A section is a contiguous area, consisting in bytes placed at consecutive positions in the file. A **hole** is also a contiguous area between two consecutive sections, but its contents is undefined and normally unused. Each section has a type. The codes associated to each section type and their interpretations are:

- - 0 (RESERVED),
 - 1 (TEXT),
 - 2 (BINARY).

Each section is considered to be a sequence of elements (records) of the same size. An element is a group of one (1) or more consecutive bytes. Elements in different sections could be of different sizes. The size of each section's elements is specified explicitly in the file's header. The name of a section (which is a string of ASCII characters) is of variable size, its actual size being specified by the "**NS**" field.

PART 1. Learning the SF format (60 min)

1. Download the following "[example.sf](#)" file, which complies perfectly the SF format.
2. We will use the following tools in order to analyze the example file. Both tools described below allow you to view the contents of a file as text (i.e. printable characters) or binary (i.e. hexadecimal codes).

Viewing a file's contents could be done by pressing F3. By default, a file contents is displayed as text, i.e. the bytes' values are interpreted as ASCII characters.

1. **on Linux:** *mc* utility (if not installed, install it by typing "sudo apt-get install mc" command in a terminal, or using "Ubunutu Software Center" application); *mc* should be run in a terminal; to switch to the hexadecimal view and back, F4 could be pressed;
2. **on Windows:** *Total Commander* (if not installed, download it from <http://www.ghisler.com/> and install it); to switch to the hexadecimal view and back, "3" could be pressed (see the "Options" menu entry).
3. Open the "*example.sf*" file in a viewer and switch to the hexadecimal view. It should look like the picture below, where you can note the file's header and data area. We illustrated the file's signatures (at both the beginning and end of the header), number of sections (8) and few section's data components.

The screenshot shows the Total Commander window with the file 'example.sf' open. The left pane displays the hexadecimal data, and the right pane displays the ASCII text. Red annotations highlight key features:

- SF SIGNATURE:** Indicated at the beginning (offset 00000000) and end (offset 00000008) of the header.
- No of sections (8):** Indicated in the header area.
- SECTION 1:** Indicated in the data area starting at offset 00000110.
- SECTION 2:** Indicated in the data area starting at offset 000001D0.
- SECTION 3:** Indicated in the data area starting at offset 00000240.

4. In the figure below we marked the header details of the first section. Please identify them on you own

The screenshot shows a hex editor window titled "Lister - [c:\Users\acolesa\Documents\os\hw1\example.sf]". The main display area shows a hex dump of the file's content. The first section header is located at offset 0x00000000 and is marked with a red box. The header contains the following fields:

- SECTION 1's SIZE (0x0064 = 100 bytes): 0x0064
- SECTION 1's TYPE (0): 0x00
- SECTION 1's METADATA: 0x00000000
- SECTION 1's ELEM SIZE (0x0005): 0x0005
- SECTION 1's NAME (SECT_RESERVED_1): 0x00000000
- SECTION 1's BEGIN OFFSET (0x0113): 0x0113
- SECTION 1's NAME SIZE (0x0F = 15): 0x000F
- SECTION 1's DATA: 0x00000000

The hex dump shows the following data for the first section header:

Offset	Hex	ASCII
00000000	12 B6 12 B6 08 00 00 00 64 00 13 01 00 00 05 00	
00000010	0F 53 45 43 54 5F 52 45 53 45 52 56 45 44 5F 31	SECT_RESERVED_1
00000020	8C 00 00 01 02 00 07 00 0A 53 45 43 54 5F 42 49	SECT_RESERVED_1
00000030	4E 5F 32 1E 00 00 02 00 00 05 00 0F 53 45 43 54	SECT_RESERVED_1
00000040	5F 52 45 53 45 52 56 45 44 5F 33 08 00 9B 02 00	SECT_RESERVED_1
00000050	00 04 00 0F 53 45 43 54 5F 52 45 53 45 52 56 45	SECT_RESERVED_1
00000060	44 5F 34 46 00 0A 02 01 00 0A 00 0B 53 45 43 54	SECT_RESERVED_1
00000070	5F 54 45 58 54 5F 35 36 00 72 03 00 00 06 00 0F	SECT_RESERVED_1
00000080	53 45 43 54 5F 52 45 53 45 52 56 45 44 5F 36 08	SECT_RESERVED_1
00000090	00 07 04 01 00 02 00 0B 53 45 43 54 5F 54 45 58	SECT_RESERVED_1
000000A0	54 5F 37 0A 00 11 04 02 00 0A 00 0A 53 45 43 54	SECT_RESERVED_1
000000B0	5F 42 49 4E 5F 38 12 B6 12 B6 00 00 00 00 00 00	SECT_RESERVED_1
000000C0	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00	SECT_RESERVED_1
000000D0	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00	SECT_RESERVED_1
000000E0	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00	SECT_RESERVED_1
000000F0	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00	SECT_RESERVED_1
00000100	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00	SECT_RESERVED_1
00000110	00 00 00 00 96 76 B8 78 02 14 2C D9 82 02 35 B8	SECT_RESERVED_1
00000120	A1 04 B7 55 15 6C 0E FE BF 22 12 92 B6 0B 60 8E	SECT_RESERVED_1
00000130	D2 6E 2B 1D 32 50 BC 49 15 E6 A0 23 65 44 6D 81	SECT_RESERVED_1
00000140	2F 0F ED A3 21 27 EA CA 41 6F 7E 4B B1 5D 96 E6	SECT_RESERVED_1
00000150	FB CB 8E C0 03 5A B5 55 62 34 2F E9 39 5F 8E 31	SECT_RESERVED_1
00000160	5D 7E DC CD AC D2 C8 17 64 69 92 14 20 7A 08 12	SECT_RESERVED_1
00000170	0E 88 80 5D 1B AD BC 00 00 00 00 00 00 00 00 00	SECT_RESERVED_1
00000180	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00	SECT_RESERVED_1
00000190	00 00 00 00 00 00 00 00 00 00 00 00 00 35 C2 9E	SECT_RESERVED_1
000001A0	93 D5 D5 96 11 91 DF CC 1E 9D DD EE D7 13 C7 7C A3	SECT_RESERVED_1

example file.

5. As an exercise, identify the header fields for another section, i.e. the fifth one, which has a TEXT type in our example. See the picture below

The fields values should be:

1. Size: 0x0046
2. Beginning offset: 0x02DA
3. Type: 0x0001 (i.e. TEXT)
4. Elements' size: 0x000A
5. Name size (NS): 0x0B
6. Name: SECT TEXT 5

PART 2. Assignment's Requirements

You are required to write a C program named "*coordinator.c*" that perform the following steps:

1.
 1. (5 min) Checks if it is given one command line parameter and validate that it corresponds to an accessible file. If not, displays an error message and terminates. We will refer from now on to that file as "CONFIG_FILE".
 2. (30 min) **Reads** one by one **all the text lines** from the CONFIG_FILE. Each line must consist in two strings of characters separated by one or more spaces. The first string corresponds to a directory path and the second one to a username. We will refer from now on to the first parameter (i.e. the directory path) as USER_DIR and to the second as USER_NAME. For each read line the program must:
 1. **validate** the line format;
 2. **validate** the USER_DIR by checking that it corresponds to an **existing, accessible** directory;
 3. **display the line**, prepended by the "VALID: " or "INVALID: " string, depending if the validation checks passed or not, respectively.
 3. (10 min) Calls a function named "authenticate()", which reads from STDIN (i.e. keyboard) a line representing a username and **checks if that username exists** in the CONFIG_FILE. If not, terminates displaying on the screen the following error message: "ERROR: invalid user". If the username exists, but its associated USER_DIR in CONFIG_FILE is invalid, then the program terminates with the error message "ERROR: Invalid user directory". Your program must restrict in the following an authenticated user's accesses only to the files located in the directory associated to that user and only if that directory is a valid one.
 4. (20 min) Calls repeatedly a function called "get_and_execute_command_line()", which reads from STDIN (i.e. keyboard) a line of characters, which correspond to a command line your program must execute. Each command line consists of a command name followed by its arguments, separated by one or more spaces, like the format illustrated below:
CMD_NAME ARG1 ARG2 ARG3 ...

The supported commands and their syntax are the following:

```
INFO file_path
SEARCH file_name [-R] [section_name section_size]
SECT_DISPLAY file_path section_name
SECT_HASH file_path section_name
EXIT
```

Your program must separate the read command line into elements, i.e. command name and its arguments, then based on the command name executes the following operations:

1. (30 min) for the INFO command, calls a function named "info(char* file_path)", which checks if the given "file_path" corresponds to an existing and valid SF file (the path is supposed to be given relative to the USER_DIR directory) and if so, displays on the screen on separate lines the following information: (1) the path of the file, (2) its number of sections and for each section, all on the same line and separated by spaces, (3) the section's name, (4) section's type, (5) section element's size, (6) section's number of elements and (7) section's size in bytes. If the path does not correspond to a SF file, the program must display the message "ERROR: not a SF file".
2. (30 min) for the SEARCH command, calls a function named "search_files_and_sections(char* file_name, char*

`section_name`, unsigned int `section_size`, unsigned char `recursive`)", which searches in the authenticated user's `USER_DIR` for all the SF files containing in their name the given "`file_name`" string, having a section whose name contains the given "`section_name`" string and a size of minimum "`section_size`" bytes. The arguments "`section_name`" and "`section_size`" could be missing, in which case any section name or section size, respectively, is accepted. If the "-R" option is specified, the search must be done in the entire file tree starting from the authenticated user's `USER_DIR`. The function must display on the screen for the matching files on different lines: (1) their path relative to the `USER_DIR` directory, (2) the matching section's name and (3) size in bytes.

3. (30 min) for the `SECT_DISPLAY` command, calls a function named "`sect_display(char* file_path, char* sect_name)`", which checks if the file corresponding to the given "`file_path`" corresponds to a SF file and the given "`sect_name`" to a TEXT section and if so, displays on the screen the contents of that section. If the checks do not succeed, the function displays on the screen the message "ERROR: not a SF file or not a TEXT section".
4. (30 min) for the `SECT_HASH` command, calls a function named "`sect_hash(char* file_path, char* sect_name)`", which checks if the file corresponding to the given "`file_path`" corresponds to a SF file and the given "`sect_name`" to a BINARY section and if so, calculate the XOR value of all its elements (taking into account the element size). Then, displays on the screen as three strings separated by a space the element size, the number of elements and the resulted XOR value as a hexadecimal value. For instance, for a binary section containing three elements of two-bytes size like "000345A5116B", the output would be "2 3 54CD". If the checks do not succeed, the function displays on the screen the message "ERROR: not a SF file or not a BINARY section".
5. (3 min) for the `EXIT` command, calls a function named "`exit()`", which terminates the program.
6. (2 min) for an unrecognized command, displays the "ERROR: Unrecognized command" message.

IMPORTANT NOTES

1. **YOU MUST UPLOAD ONLY YOUR "coordinator.c" file!**
2. **You program must successfully compile with a command line like "`gcc -Wall -Werror coordinator.c -o coordinator`" to be further evaluated.** The "-Wall" option says the compiler to report anything that could be wrong from its point of view, while the "-Werror" says it to report any warning as an error and do not accept it. As a consequence, your program must compile without any error and warning at all in order to be admitted for evaluation.
3. You are required and **restricted to use only the OS system calls**, i.e. low-level functions, not higher-level one, in your entire solution, in all lab assignments. For instance, regarding the file accesses, you **MUST use system calls** like `open()`, `read()`, `write()` etc., but NOT higher-level functions like `fopen()`, `fgets()`, `fscanf()`, `fprintf()` etc. The only accepted exceptions from this requirement are the functions to read from STDIN or display to STDOUT / STDERR, like `scanf()`, `printf()`, `perror()` and functions for string manipulation and conversion like `sscanf()`, `snprintf()`.
4. For testing purposes you can download the following [archive](#), containing a directory with few testing SF files. Please note that there could be both valid and invalid SF format files in the given archive. Start by reading the README file in the given archive.

5. A **SF format validation must check for both file signatures and the header structure integrity**, i.e. checks if the number of meta-data areas corresponds to the reported number of sections.
6. It could be supposed that any SF file contains only sections with distinct name, having no two sections with the same name. Though, there could be different sections whose name contain a common substring (like, for instance, "SECT_1", "SECT_2", "MY_SECT").
7. The text sections were generated using the Linux convention, so contain just one byte (value 10) for the new-line ('\n') separator. Take care that the last line could end without a '\n', i.e. the last character in the file could be different by '\n'.
8. For string tokenization (i.e. separate a string into elements based on specific separators, like spaces) we recommend you using the `strtok()` or `strtok_r()` functions.
9. You can automate your program runs by redirecting its STDIN to a text file prepared to contain the commands normally given from the keyboard. This way you can easily and efficiently run the same test more times for debugging purposes. This is actually the way we will test your solutions.
10. Extra credit (1 point) will be given to you if you use regular expressions to specify string patterns for file and section names for the SEARCH command (see manual page REGEX(3) for details).